

Python Basics - Individual

● Graded

Student

Anand Patel

Total Points

84.5 / 100 pts

Autograder Score

3.0 / 5.0

Failed Tests

test_data_dir (test_files.TestFiles) (0/1)

test_list_files (test_files.TestFiles) (2/3)

Passed Tests

test_zip (test_files.TestFiles) (1/1)

Question 2

1.2 - for loop

1 / 1 pt

✓ - 0 pts Correct

Question 3

1.3 - verify solution

1 / 1 pt

✓ - 0 pts Correct

Question 4

1.4 - efficiency check

1 / 1 pt

✓ - 0 pts Correct

Question 5

1.5 - discussion

0.5 / 1 pt

✓ - 0.5 pts Missing discussion of design details.

Question 6

1.6 - design, implement, verify (less efficient)

2 / 3 pts

✓ - 1 pt Verification unclear.

Question 7

1.7 - discussion (less efficient)

1 / 1 pt

✓ - 0 pts Good discussion.

Question 8

1.8 - Comparing Implementations

1 / 2 pts

✓ - 1 pt Minimal explanation provided.

Question 9

2.0.1 - Import

2 / 2 pts

✓ - 0 pts Correct

Question 10

2.0.2 - Load

1 / 1 pt

✓ - 0 pts Correct

Question 11

2.0.3 - Display data

1 / 1 pt

✓ - 0 pts Correct

Question 12

2.0.4 - Display data information

1 / 1 pt

✓ - 0 pts Correct

Question 13

2.1.1 - Remove columns

1 / 2 pts

✓ - 1 pt No statement

Question 14

2.1.2 - Rename columns

1 / 2 pts

✓ - 1 pt No statement

Question 15

2.1.3 - Remove duplicated rows

3 / 3 pts

✓ - 0 pts Correct

Question 16

2.1.4 - Null values

2 / 2 pts

✓ - 0 pts Correct

Question 17

2.1.5 - Handling missing values

3 / 3 pts

✓ - 0 pts Correct

Question 18

2.1.6 - Reasoning

0 / 3 pts

✓ - 3 pts have not answered the question

Question 19

2.2.1 - Add HP_Type

2 / 2 pts

✓ - 0 pts Correct

Question 20

2.2.2 - Add Price_class

2 / 2 pts

✓ - 0 pts Correct

Question 21

2.2.3 - Save modified data

1 / 1 pt

✓ - 0 pts Correct

Question 22

2.3.1 - Mean price

1 / 1 pt

✓ - 0 pts Correct

Question 23

2.3.2 - Count cars by year

1 / 1 pt

✓ - 0 pts Correct

Question 24

2.3.3 - Unique Makes

1 / 1 pt

✓ - 0 pts Correct

Question 25

2.3.4 - Fuel efficiency

1.5 / 2 pts

✓ - 0.5 pts Report the number of unique vehicle makers.

Question 26

2.3.5 - Cross-tabulation

2 / 2 pts

✓ - 0 pts Correct

Question 27

2.3.6 - Isolating specific features

1 / 1 pt

✓ - 0 pts Correct

Question 28

2.3.7 - Isolating specific features

2 / 2 pts

✓ - 0 pts Correct

Question 29

2.3.8 - Counting specific features

2 / 2 pts

✓ - 0 pts Correct

Question 30

2.3.9 - Grouping data

1.5 / 2 pts

✓ - 0.5 pts You should display the values as integers

Question 31

2.3.10 - Reporting extreme values

1 / 1 pt

✓ - 0 pts Correct

Question 32

3.1 - Prepare data

5 / 5 pts

✓ - 0 pts Correct formulations.

Question 33

3.2 - Define a Tournament class

26 / 30 pts

✓ - 2 pts A more optimal selection of cars could be guaranteed found by implementation of the classical 0/1 knapsack problem (as opposed to fractional or greedy solution).

✓ - 1 pt Winners should be permitted to `_purchase_inventory` again between matches.

✓ - 1 pt Missing full set of complete docstrings.

Question 34

3.3 - Execute your Tournament class

5 / 5 pts

✓ - 0 pts Good execution demonstration.

Question 35

3.4 - Compare the results of multiple Tournament executions

4 / 5 pts

✓ + 4 pts This could be written more generally, to operate on any number of Tournaments. You might even overload a comparison operator to compare instances of Tournaments.

Autograder Results

test_data_dir (test_files.TestFiles) (0/1)

Test Failed: False is not true : Missing data directory.

test_list_files (test_files.TestFiles) (2/3)

Found *CS2PP22_CW1.ipynb
Missing *CS2PP22_CW1.html
Found *cardata_modified.csv

test_zip (test_files.TestFiles) (1/1)

Submitted Files

Front page of the student's submission (the following are compulsory):

Module Code: CS2PP22

Assignment report Title: CW1 Individual


Date (when the work completed): 16/02/2024

Actual hrs spent for the assignment: 20

CS2PP22 Programming in Python for Data Science**Instructions:**


- Write Python code to perform each of the following sub-tasks.
 - Try to follow the PEP 8 – Style Guide for Python Code: <https://peps.python.org/pep-0008>
 - Function and variable type annotations are not required.
- Some parts of this assignment may require further self-study of Python documentations or other resources.
- You may also refer to other documentation/self-study resources, such as those suggested in the Lecture Notes or a multitude of other resources
- Blank code and markdown cells are provided for each sub-task, however you will likely need to create additional cells to provide further explanation

Items to be submitted: 1. A modified version of this Jupyter notebook file (`.ipynb`) - This is to be submitted already **fully executed** i4. `cardata_modified.csv`6. * *Optional* * : Any functions and classes created for this Task may be written in a separate `.py` module file and `import` ed to this space. The**10 marks**

- A **network** or **graph**, G , consists of a set of **nodes** (or vertices), V , and edges, E .
 - An edge is a pair of nodes (a,b) denoting the nodes connected by the edge. $G = (V,E)$ 
 - (images/CompareNetworks.png) - Networks can be represented in various different structures.
- We can represent networks using a simple list of all the nodes and all the edges in the network:

$$V = a, b, c, d, e$$



- Another data structure that we could use is to build a **neighbour list** for every node.
 - For every node, x , we maintain a list of all the neighbours, y . 

Data Reference:

The file `data/dolphins.tsv` contains a representation of a social network dataset where dolphins have links between them if they frequently associated with one another.

- Taken from the Koblenz Network Collection by the University of Koblenz–Landau. Dolphins network dataset – KONECT, April 2017. <http://konect.cc/networks/dolphins/>.

To read in the *tab separated* network file, we need to read each line in the file. To do this we just treat the file object as an iterator.

- This is an example of using a context manager: https://book.pythontips.com/en/latest/context_managers.html

- Here `line` will be a string for each line in the file. The `.split(x)` method splits a string into a list of substrings for each occurrence of the character `x` (in this case the tab: `'\t'`). This code reads the file and places the data in an edge list representation: ``python # Create an empty edge set edges = set() with open('data/dolphins.tsv', 'r') as file: for line in file: a, b = line.split('\t') e = (int(a), int(b)) edges.add(e) `` ---

****Instructions:**** 1. Begin with the resulting `edges` variable. 2. Use a `for` loop to populate a `dict` that will contain the neighbour list network representation. 3. ****Verify**** that your solution matches that for nodes 55, 2, and 20 (shown below). 4. With your entire neighbour list code in a single Jupyter notebook cell, use the `timeit` Magic command to report the performance of your code, executing 19 runs of 2000 loops each. - There should be no network data output from this cell. 5. Discuss your solution, describing the Python constructs/tools you have used, and the design of your implementation. 6. Design, implement, and verify the results of a second, less efficient implementation and record its equivalent performance metrics as above. 7. Discuss your second solution, describing the Python constructs/tools you have used, and the design of your less efficient implementation. 8. Explain why the efficiency differs between your two implementations. --- There will be many nodes in the final dataset. The result should have the form: ``python {55: {2, 7, 8, 14, 20, 42, 58}, 2: {18, 20, 27, 28, 29, 37, 42, 55}, 20: {2, 8, 31, 55},... `` This shows that dolphin `55` is frequently associated with dolphins `2, 7, 8, 14, 20, 42, and 58`.

In [8]:

```
# Create an empty edge set
edges = set()

with open('dolphins.tsv', 'r') as file:
    for line in file:
        a, b = line.split('\t')
        e = (int(a), int(b))
        edges.add(e)
```

In [31]:

```
%%timeit -r 19 -n 2000
network = {}
for edge in edges:
    a, b = edge
    if a not in network:
        network[a] = set()
    if b not in network:
        network[b] = set()
    network[a].add(b)
    network[b].add(a)
```

59.6 μ s \pm 3.96 μ s per loop (mean \pm std. dev. of 19 runs, 2000 loops each)

In [36]:

```
# Verifying my more efficient solution
nodes = [55, 2, 20]
for node in nodes:
    print(node, network[node])
```

```
55 {2, 7, 8, 42, 14, 20, 58}
2 {37, 42, 18, 20, 55, 27, 28, 29}
20 {8, 2, 31, 55}
```

In [32]:

Discussion: This approach takes advantage of the speed of dictionaries and the uniqueness property of sets so that each neighbor is listed once for each node. This prevents redundant/unnecessary checks.

In [38]:

```
%%timeit -r 19 -n 2000
network = {}
for a, b in edges:
    if isinstance(network.get(a), set):
        network[a].add(b)
    else:
        network[a] = {b}
    if isinstance(network.get(b), set):
        network[b].add(a)
```

```
else:  
    network[b] = {a}
```

83.3 μ s \pm 3.73 μ s per loop (mean \pm std. dev. of 19 runs, 2000 loops each)

In []:

Discussion: dict.get is less efficient than the first solution as it will always return a value as it does a check if the key exists within the set and isinstance does a type check of the value. These 2 extra steps for each edge will increase the overhead compared to the first solution

Data Reference:

- Car Features and MSRP Data: `cardata.csv`

—This dataset includes car features such as make, model, year, and engine type, as scraped from Edmunds and Twitter. It is often used to

- Source: <https://www.kaggle.com/datasets/CooperUnion/cardataset>

—Each *row* corresponds to a single kind of vehicle.

- The **columns** correspond to:

Column	Description
Make	Car maker
Model	Car model
Year	Car year (Marketing)
Engine Fuel Type	Type of engine fuel category
Engine HP	Engine horsepower (HP)
Engine Cylinders	Number of engine cylinders
Transmission Type	Type of transmission category
Driven_Wheels	Drive wheel category
Number of Doors	Number of doors
Market Category	Market category
Vehicle Size	Vehicle size category
Vehicle Style	Vehicle style category
highway MPG	Highway fuel efficiency in miles per gallon

Column	Description
city mpg	City fuel efficiency in miles per gallon
Popularity	Twitter-based popularity metric
MSRP	Manufacturer suggested retail price (USD)

2.0. Analysis Preparation

5 Marks

2.0.1. Import [2 marks](#)

Import the libraries that will be used in your solution.

In [3]:

```
import numpy as np
import pandas as pd
```

2.0.2. Load data [1 mark](#)

Locate the data file `cardata.csv` within the zipped file you have downloaded from Blackboard (under `./data/Task1/`) and read the data from file to a `pandas` `DataFrame`.

In [4]:

```
cars = pd.read_csv('cardata.csv')
```

2.0.3. Display data [1 mark](#)

Display the first 5 rows of the `DataFrame`.

In [5]:

```
cars.head()
```

```
Out [5]:
```

	Make	Model	Year	Engine	Fuel Type	Engine HP \
0	BMW	1 Series M	2011	premium	unleaded (required)	335.0
1	BMW	1 Series	2011	premium	unleaded (required)	300.0
2	BMW	1 Series	2011	premium	unleaded (required)	300.0
3	BMW	1 Series	2011	premium	unleaded (required)	230.0
4	BMW	1 Series	2011	premium	unleaded (required)	230.0

	Engine Cylinders	Transmission	Type	Driven_Wheels	Number of Doors \
0	6.0	MANUAL	rear wheel drive	2.0	
1	6.0	MANUAL	rear wheel drive	2.0	
2	6.0	MANUAL	rear wheel drive	2.0	
3	6.0	MANUAL	rear wheel drive	2.0	
4	6.0	MANUAL	rear wheel drive	2.0	

	Market Category	Vehicle Size	Vehicle Style \
0	Factory Tuner,Luxury,High-Performance	Compact	Coupe
1	Luxury,Performance	Compact	Convertible
2	Luxury,High-Performance	Compact	Coupe
3	Luxury,Performance	Compact	Coupe
4	Luxury	Compact	Convertible

	highway MPG	city mpg	Popularity	MSRP
0	26	19	3916	46135
1	28	19	3916	40650
2	28	20	3916	36350
3	28	18	3916	29450
4	28	18	3916	34500

2.0.4. Display data information [1 mark](#)

In one or more code cells, display the following information of the `DataFrame`:

- number of rows
- number of columns
- column names and data types

Following the display of this information, use a Markdown cell to write:

- one sentence summarising the numbers of rows and columns,
- and one sentence describing the unique data types.

In [6]:

```
cars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11914 entries, 0 to 11913
Data columns (total 16 columns):
#   Column              Non-Null Count  Dtype  
---  -
0   Make                 11914 non-null object  
1   Model                11914 non-null object  
2   Year                 11914 non-null int64  
3   Engine Fuel Type     11911 non-null object  
4   Engine HP            11845 non-null float64  
5   Engine Cylinders     11884 non-null float64  
6   Transmission Type    11914 non-null object  
7   Driven_Wheels        11914 non-null object  
8   Number of Doors      11908 non-null float64  
9   Market Category      8172 non-null  object  
10  Vehicle Size         11914 non-null object  
11  Vehicle Style        11914 non-null object  
12  highway MPG          11914 non-null int64  
13  city mpg             11914 non-null int64  
14  Popularity           11914 non-null int64  
15  MSRP                 11914 non-null int64  
dtypes: float64(3), int64(5), object(8)
memory usage: 1.5+ MB
```

This `DataFrame` contains 11,914 rows and 16 columns.

The data types present include strings (recorded as `object`), 64-bit integers (`int64`), and 64-bit floating point values (`float64`).

2.1. Data Cleaning

[15 Marks](#)

2.1.1. Remove columns [2 marks](#)

Drop the following columns from the `DataFrame`:

- Engine Fuel Type
- Market Category
- Number of Doors
- Vehicle Size

Then, **display** the resulting `DataFrame` and **write a statement**, explaining why one might choose to exclude these features.

In [7]:

```
cars = cars.drop(columns=['Engine Fuel Type', 'Market Category', 'Number of Doors', 'Vehicle Size'])
cars
```

Out [7]:

```
   Make  Model Year  Engine HP  Engine Cylinders \
0  BMW  1 Series M  2011    335.0             6.0
1  BMW  1 Series  2011    300.0             6.0
2  BMW  1 Series  2011    300.0             6.0
3  BMW  1 Series  2011    230.0             6.0
```



```

4      BMW  1 Series 2011   230.0      6.0
...    ...    ...    ...    ...
11909  Acura   ZDX 2012   300.0      6.0
11910  Acura   ZDX 2012   300.0      6.0
11911  Acura   ZDX 2012   300.0      6.0
11912  Acura   ZDX 2013   300.0      6.0
11913  Lincoln Zephyr 2006   221.0      6.0

      Transmission Type   Driven_Wheels Vehicle Style highway MPG \
0          MANUAL    rear wheel drive      Coupe      26
1          MANUAL    rear wheel drive  Convertible      28
2          MANUAL    rear wheel drive      Coupe      28
3          MANUAL    rear wheel drive      Coupe      28
4          MANUAL    rear wheel drive  Convertible      28
...    ...    ...    ...    ...
11909  AUTOMATIC    all wheel drive 4dr Hatchback      23
11910  AUTOMATIC    all wheel drive 4dr Hatchback      23
11911  AUTOMATIC    all wheel drive 4dr Hatchback      23
11912  AUTOMATIC    all wheel drive 4dr Hatchback      23
11913  AUTOMATIC front wheel drive      Sedan      26

      city mpg Popularity  MSRP
0         19     3916 46135
1         19     3916 40650
2         20     3916 36350
3         18     3916 29450
4         18     3916 34500
...    ...    ...    ...
11909     16      204 46120
11910     16      204 56670
11911     16      204 50620
11912     16      204 50920
11913     17       61 28995

```

[11914 rows x 12 columns]

2.1.2. Rename columns [2 marks](#)

Rename the following columns:

Column	New name
Engine HP	HP
Engine Cylinders	Cylinders
Transmission Type	Transmission
Driven_Wheels	Drive Mode
highway MPG	MPG-H
city mpg	MPG-C
MSRP	Price

Then, **display** the resulting `DataFrame` and **write a statement**, discussing why one might change the names in this way.

In [8]:

```
cars = cars.rename(columns = {'Engine HP':'HP','Engine Cylinders':'Cylinders','Transmission Type':'Transmission','Driven_Wheels':'Drive Mode','highway MPG':'MPG-H','city mpg':'MPG-C','MSRP':'Price'})
cars
```

Out [8]:

	Make	Model	Year	HP	Cylinders	Transmission	\
0	BMW	1 Series M	2011	335.0	6.0	MANUAL	
1	BMW	1 Series	2011	300.0	6.0	MANUAL	
2	BMW	1 Series	2011	300.0	6.0	MANUAL	
3	BMW	1 Series	2011	230.0	6.0	MANUAL	
4	BMW	1 Series	2011	230.0	6.0	MANUAL	
...
11909	Acura	ZDX	2012	300.0	6.0	AUTOMATIC	
11910	Acura	ZDX	2012	300.0	6.0	AUTOMATIC	
11911	Acura	ZDX	2012	300.0	6.0	AUTOMATIC	
11912	Acura	ZDX	2013	300.0	6.0	AUTOMATIC	
11913	Lincoln	Zephyr	2006	221.0	6.0	AUTOMATIC	

	Drive Mode	Vehicle Style	MPG-H	MPG-C	Popularity	Price
0	rear wheel drive	Coupe	26	19	3916	46135
1	rear wheel drive	Convertible	28	19	3916	40650
2	rear wheel drive	Coupe	28	20	3916	36350
3	rear wheel drive	Coupe	28	18	3916	29450
4	rear wheel drive	Convertible	28	18	3916	34500
...
11909	all wheel drive	4dr Hatchback	23	16	204	46120
11910	all wheel drive	4dr Hatchback	23	16	204	56670
11911	all wheel drive	4dr Hatchback	23	16	204	50620
11912	all wheel drive	4dr Hatchback	23	16	204	50920
11913	front wheel drive	Sedan	26	17	61	28995

[11914 rows x 12 columns]

2.1.3. Remove duplicated rows [3 marks](#)

Use an `f-string` to `print()` the number of irrelevant duplicate rows (count repeats only, not both originals and repeats).

Drop the duplicated rows, retain the resulting `DataFrame`, and show the resulting number of rows in the new `DataFrame`.

```
In [9]: print(f'Duplicate rows: {cars.duplicated().sum()}')
cars = cars.drop_duplicates()
print(f'Rows after dropping duplicates: {len(cars)}')
```

Duplicate rows: 803
Rows after dropping duplicates: 11111

2.1.4. Null values [2 marks](#)

Report the number of null values remaining in each column.

```
In [10]: null_counts = cars.isnull().sum()

print(null_counts)
```

```
Make      0
Model     0
Year      0
HP        69
Cylinders 30
Transmission  0
Drive Mode  0
Vehicle Style  0
MPG-H     0
MPG-C     0
Popularity  0
Price     0
dtype: int64
```

2.1.5. Handling missing values [3 marks](#)

- For the `HP` column, replace missing values with the column mean.
- Drop all other rows that still contain missing values.
- Display the final number of rows in the dataframe.

```
In [11]: hp_mean = np.mean(cars.HP)
cars = cars.fillna(value={"HP": hp_mean})
cars.isnull().sum()
```

```
Out [11]: Make      0
Model    0
Year     0
HP       0
Cylinders 30
Transmission 0
Drive Mode 0
Vehicle Style 0
MPG-H     0
MPG-C     0
Popularity 0
Price     0
dtype: int64
```

```
In [12]: cars = cars.dropna()

print('Rows: ', cars.shape[0])
```

Rows: 11081

2.1.6. Reasoning [3 marks](#)

In the previous step, we eliminated entries with missing data from the dataset. If this dataset were to be used to train a machine learning model to predict vehicle prices, what is a potential drawback of this approach? Describe an alternative approach and any related caveats of which we should be aware.

2.2. Creating New Columns

[5 Marks](#)

2.2.1. Add `HP_Type` [2 marks](#)

Using an implementation of **list comprehension**, create a new column, `HP_Type`, such that,

- if a car's `HP` is greater than or equal to 300, `HP_Type` is 'high'.
- Otherwise, `HP_Type` is 'low'.

Then, display the resulting `DataFrame`.

```
In [13]: cars['HP_Type'] = ['high' if hp >= 300 else 'low' for hp in cars['HP']]
cars
```

```
Out [13]:   Make  Model Year  HP  Cylinders Transmission \
0   BMW  1 Series M  2011  335.0    6.0    MANUAL
1   BMW  1 Series  2011  300.0    6.0    MANUAL
2   BMW  1 Series  2011  300.0    6.0    MANUAL
3   BMW  1 Series  2011  230.0    6.0    MANUAL
4   BMW  1 Series  2011  230.0    6.0    MANUAL
...   ...   ...   ...   ...   ...   ...
11909  Acura    ZDX  2012  300.0    6.0  AUTOMATIC
11910  Acura    ZDX  2012  300.0    6.0  AUTOMATIC
11911  Acura    ZDX  2012  300.0    6.0  AUTOMATIC
11912  Acura    ZDX  2013  300.0    6.0  AUTOMATIC
```

```

11913 Lincoln    Zephyr 2006 221.0    6.0  AUTOMATIC

      Drive Mode Vehicle Style MPG-H MPG-C Popularity Price \
0   rear wheel drive      Coupe  26   19   3916 46135
1   rear wheel drive  Convertible  28   19   3916 40650
2   rear wheel drive      Coupe  28   20   3916 36350
3   rear wheel drive      Coupe  28   18   3916 29450
4   rear wheel drive  Convertible  28   18   3916 34500
...
11909 all wheel drive 4dr Hatchback 23   16   204 46120
11910 all wheel drive 4dr Hatchback 23   16   204 56670
11911 all wheel drive 4dr Hatchback 23   16   204 50620
11912 all wheel drive 4dr Hatchback 23   16   204 50920
11913 front wheel drive      Sedan  26   17   61 28995

```

```

      HP_Type
0   high
1   high
2   high
3   low
4   low
...
11909 high
11910 high
11911 high
11912 high
11913 low

```

```
[11081 rows x 13 columns]
```

2.2.2. Add `Price_class` [2 marks](#)

Using an implementation of a **function**, create a new column, `Price_class`, such that it becomes equal to:

- 'high', if `Price` is greater than or equal to 50,000
- 'mid', if `Price` is between 30,000 (inclusive) and 50,000 (exclusive), and
- 'low', if `Price` is below 30,000.

Then, display the resulting `DataFrame`.

In [14]:

```

def price_category(price):
    if price >= 50000:
        return 'high'
    elif price >= 30000:
        return 'mid'
    else:
        return 'low'

cars['Price_class'] = cars['Price'].apply(price_category)
cars

```

Out [14]:

```

      Make  Model Year  HP  Cylinders Transmission \
0   BMW  1 Series M 2011 335.0    6.0    MANUAL
1   BMW  1 Series 2011 300.0    6.0    MANUAL
2   BMW  1 Series 2011 300.0    6.0    MANUAL
3   BMW  1 Series 2011 230.0    6.0    MANUAL
4   BMW  1 Series 2011 230.0    6.0    MANUAL
...
11909 Acura    ZDX 2012 300.0    6.0  AUTOMATIC
11910 Acura    ZDX 2012 300.0    6.0  AUTOMATIC
11911 Acura    ZDX 2012 300.0    6.0  AUTOMATIC
11912 Acura    ZDX 2013 300.0    6.0  AUTOMATIC
11913 Lincoln  Zephyr 2006 221.0    6.0  AUTOMATIC

      Drive Mode Vehicle Style MPG-H MPG-C Popularity Price \
0   rear wheel drive      Coupe  26   19   3916 46135
1   rear wheel drive  Convertible  28   19   3916 40650
2   rear wheel drive      Coupe  28   20   3916 36350
3   rear wheel drive      Coupe  28   18   3916 29450
4   rear wheel drive  Convertible  28   18   3916 34500

```

```

...      ...      ...      ...      ...
11909  all wheel drive  4dr Hatchback  23  16      204 46120
11910  all wheel drive  4dr Hatchback  23  16      204 56670
11911  all wheel drive  4dr Hatchback  23  16      204 50620
11912  all wheel drive  4dr Hatchback  23  16      204 50920
11913  front wheel drive      Sedan    26  17      61 28995

```

```

HP_Type Price_class
0    high      mid
1    high      mid
2    high      mid
3    low       low
4    low      mid

...      ...      ...
11909  high      mid
11910  high      high
11911  high      high
11912  high      high
11913  low       low

```

[11081 rows x 14 columns]

2.2.3 Save modified data [1 mark](#)

Save the modified `DataFrame` to a new comma-separated file called `cardata_modified.csv` to be stored under the `data/Task1/` directory.

Do not include the row indices in the file.

```
In [15]: cars.to_csv('cardata_modified.csv', index=False)
```

2.3. Exploratory Data Analysis

[15 Marks](#)

2.3.1. Mean price [1 mark](#)

Find the mean `Price` of all vehicles. Report the solution rounded to 2 decimal places.

```
In [16]: round(cars['Price'].mean(), 2)
```

```
Out [16]: 41957.25
```

2.3.2. Count cars by year [1 mark](#)

Report the number of cars in each `Year`.

```
In [17]: cars['Year'].value_counts()
```

```
Out [17]: 2015    2048
2016    2046
2017    1598
2014     542
2009     356
2012     350
2007     334
2013     325
2008     322
2011     279
2010     276
2003     238
2004     235
2005     213
2002     205
```

```
2006 194
2001 168
1997 162
1993 159
1998 144
1992 127
1994 125
2000 115
1995 114
1999 114
1996 113
1991 102
1990 77
Name: Year, dtype: int64
```

2.3.3. Unique Makes [1 mark](#)

Report a list of unique values of `Make`, sorted in ascending alphabetical order.

```
In [18]: sorted(list(set(cars.Make)))
```

```
Out [18]: ['Acura',
'Afa Romeo',
'Aston Martin',
'Audi',
'BMW',
'Bentley',
'Bugatti',
'Buick',
'Cadillac',
'Chevrolet',
'Chrysler',
'Dodge',
'FIAT',
'Ferrari',
'Ford',
'GMC',
'Genesis',
'HUMMER',
'Honda',
'Hyundai',
'Infiniti',
'Kia',
'Lamborghini',
'Land Rover',
'Lexus',
'Lincoln',
'Lotus',
'Maserati',
'Maybach',
'Mazda',
'McLaren',
'Mercedes-Benz',
'Mitsubishi',
'Nissan',
'Oldsmobile',
'Plymouth',
'Pontiac',
'Porsche',
'Rolls-Royce',
'Saab',
'Scion',
'Spyker',
'Subaru',
'Suzuki',
'Tesla',
'Toyota',
'Volkswagen',
'Volvo']
```

2.3.4. Fuel efficiency [2 marks](#)

Report the number of unique vehicle makers.

Display the mean `MPG-H` and mean `MPG-C` for each `Model` of each `Make` in a `MultiIndex DataFrame`. Report the values rounded to 2 decimal places.

Verify that the expected number of vehicle makers are represented in the resulting `DataFrame`.

```
In [19]: cars.Make.unique().size
mean_mpg = cars.groupby(['Make', 'Model'])[['MPG-H', 'MPG-C']].mean().round(2)
mean_mpg
```

```
Out [19]:      MPG-H  MPG-C
Make Model
Acura CL    26.78  17.00
   ILX     35.12  24.62
   ILX Hybrid 38.00  39.00
   Integra  28.17  21.83
   Legend   23.19  16.00
...
Volvo V90    23.00  16.00
   XC       23.00  17.00
   XC60     27.13  19.87
   XC70     27.50  20.17
   XC90     25.29  20.47

[923 rows x 2 columns]
```

2.3.5. Cross-tabulation [2 marks](#)

Display in a single `DataFrame` the number of car entries for each `Drive Mode` and `Transmission` combination, with a `Total` column (showing the sum of the rows) and a `Total` row (showing the sum of the columns) in the `margins`.

```
In [20]: pd.crosstab(index=cars['Drive Mode'], columns=cars['Transmission'], margins=True, margins_name='Total')
```

```
Out [20]: Transmission  AUTOMATED_MANUAL  AUTOMATIC  DIRECT_DRIVE  MANUAL  UNKNOWN \
Drive Mode
all wheel drive      198      1897      11    202      0
four wheel drive       0       979       0    291      2
front wheel drive     231     2861      36   1209      4
rear wheel drive      124     2101      11    918      6
Total                553     7838      58   2620     12

Transmission  Total
Drive Mode
all wheel drive  2308
four wheel drive 1272
front wheel drive 4341
rear wheel drive 3160
Total          11081
```

2.3.6. Isolating specific features [1 mark](#)

Display data corresponding to the `Year` 2017 for Porsche's Macan model.

```
In [21]: cars[(cars['Year'] == 2017) & (cars['Make'] == 'Porsche') & (cars['Model'] == 'Macan')]
```

```
Out [21]:   Make Model Year  HP Cylinders  Transmission \
6630 Porsche Macan  2017  360.0      6.0  AUTOMATED_MANUAL
6631 Porsche Macan  2017  340.0      6.0  AUTOMATED_MANUAL
6632 Porsche Macan  2017  400.0      6.0  AUTOMATED_MANUAL
6633 Porsche Macan  2017  252.0      4.0  AUTOMATED_MANUAL
```

	Drive Mode	Vehicle Style	MPG-H	MPG-C	Popularity	Price	HP_Type	\
6630	all wheel drive	4dr SUV	23	17	1715	67200	high	
6631	all wheel drive	4dr SUV	23	17	1715	54400	high	
6632	all wheel drive	4dr SUV	23	17	1715	76000	high	
6633	all wheel drive	4dr SUV	25	20	1715	47500	low	

	Price_class
6630	high
6631	high
6632	high
6633	mid

2.3.7. Isolating specific features [2 marks](#)

Display data corresponding to the `Year` 2015 for makers FIAT and Scion that have entries with automatic `Transmission` in a **single** `DataFrame`.

In [22]: `cars[(cars['Year'] == 2015) & ((cars['Make'] == 'FIAT') | (cars['Make'] == 'Scion')) & (cars.Transmission == 'AUTOMATIC')]`

Out [22]:

	Make	Model	Year	HP	Cylinders	Transmission	Drive Mode	\
572	FIAT	500L	2015	160.0	4.0	AUTOMATIC	front wheel drive	
4897	Scion	FR-S	2015	200.0	4.0	AUTOMATIC	rear wheel drive	
4898	Scion	FR-S	2015	200.0	4.0	AUTOMATIC	rear wheel drive	
5962	Scion	iQ	2015	94.0	4.0	AUTOMATIC	front wheel drive	
10312	Scion	tC	2015	179.0	4.0	AUTOMATIC	front wheel drive	
10313	Scion	tC	2015	179.0	4.0	AUTOMATIC	front wheel drive	
11548	Scion	xB	2015	158.0	4.0	AUTOMATIC	front wheel drive	
11549	Scion	xB	2015	158.0	4.0	AUTOMATIC	front wheel drive	

	Vehicle Style	MPG-H	MPG-C	Popularity	Price	HP_Type	Price_class
572	Wagon	30	22	819	24695	low	low
4897	Coupe	34	25	105	31090	low	mid
4898	Coupe	34	25	105	26000	low	low
5962	2dr Hatchback	37	36	105	15665	low	low
10312	2dr Hatchback	31	23	105	20360	low	low
10313	2dr Hatchback	31	23	105	24340	low	low
11548	Wagon	28	22	105	18070	low	low
11549	Wagon	28	22	105	19685	low	low

2.3.8. Counting specific features [2 marks](#)

For BMWs with `Price` greater than 30,000, report the number of cars for each `Transmission` category.

In [23]: `cars[(cars['Make'] == 'BMW') & (cars['Price'] > 30000)].groupby('Transmission').size()`

Out [23]:

Transmission	
AUTOMATED_MANUAL	18
AUTOMATIC	240
DIRECT_DRIVE	4
MANUAL	51

dtype: int64

2.3.9. Grouping data [2 marks](#)

For 2017 data, show the minimum and maximum `Price`, as well as minimum and maximum `HP` for each `Make`.

Display the values as **integers** in a **single** `DataFrame`.

In [24]: `cars[(cars['Year'] == 2017)].groupby(['Make'])[['Price', 'HP']].agg([np.min, np.max])`

Out [24]:

	Price	HP		
	min	max	min	max
Make				

Acura	27990	156000	201.0	573.000000
Audi	31200	189900	186.0	610.000000
BMW	33100	137000	170.0	600.000000
Buick	21065	49625	138.0	310.000000
Cadillac	34595	97795	265.0	640.000000
Chevrolet	13000	92395	98.0	650.000000
Chrysler	21995	45270	184.0	300.000000
Dodge	20995	118795	173.0	707.000000
FIAT	15990	31800	101.0	252.588752
Ford	14130	68996	120.0	526.000000
GMC	24070	71665	182.0	420.000000
Genesis	41400	54550	311.0	420.000000
Honda	15990	47070	130.0	280.000000
Hyundai	17150	41150	128.0	293.000000
Infiniti	33950	53300	208.0	400.000000
Kia	14165	45700	138.0	290.000000
Land Rover	37695	62500	240.0	240.000000
Lexus	31250	89380	134.0	467.000000
Lincoln	32720	76645	240.0	380.000000
Lotus	91900	91900	400.0	400.000000
Maserati	71600	145500	345.0	523.000000
Mazda	17845	30695	146.0	184.000000
Mercedes-Benz	32400	247900	177.0	621.000000
Mitsubishi	12995	31695	78.0	224.000000
Nissan	11990	109990	109.0	565.000000
Porsche	47500	200400	252.0	580.000000
Subaru	18395	39995	148.0	305.000000
Toyota	15250	84325	106.0	381.000000
Volkswagen	17895	60195	150.0	292.000000
Volvo	33950	57200	240.0	316.000000

2.3.10. Reporting extreme values [1 mark](#)


Determine which `Make` has the highest mean `Price`. Display this `Make` and its corresponding mean `Price` (rounded to 2 decimal places).

In [25]:

```
cars.groupby('Make')['Price'].mean().sort_values().tail(1).round(2)
```

Out [25]:
 Make
 Bugatti 1757223.67
 Name: Price, dtype: float64

[40 Marks](#)



Design a `Tournament` class to represent **teams** of car efficiency enthusiasts that can collect an **inventory** of fuel efficient cars by finding sponsorship from car makers.

Once **sponsors** are arranged and teams are formed, teams will **compete** by facing off head-to-head (pairwise, bracket-style) until only one `Tournament.champion` prevails.

Competitors perform by driving each car in their inventory the distance that **1 gallon** of fuel will permit them to travel in **highway** situations.

The **winning team** of a match will have collectively driven their vehicles the furthest. That is, the sum of the `MPG-H` ratings in their inventory will be the **score** that decides the winner.

3.1 Prepare data [5 marks](#)

Begin with the data you have saved in the `cardata_modified.csv` file.

Read in and modify this data so that:

- The `Price`s are rounded to the **nearest \$100**.
- Only cars made **after the Year 2000** are retained in the `DataFrame`.
- Only car `Make`s with **more than 55 entries** in the dataset are retained in the `DataFrame`.

Teams will be able to choose cars for their inventory from this modified `DataFrame`.

In [26]:

```
cars = pd.read_csv('cardata_modified.csv')
cars['Price'] = cars['Price'].round(-2)
cars = cars[cars['Year'] > 2000]

cars = cars[cars.groupby('Make').Make.transform('count') > 55]
cars.Make.value_counts()
```

Out [26]:

```
Chevrolet    959
Ford         693
Toyota       576
Volkswagen   545
Nissan        486
GMC          427
Dodge        414
Honda        410
Cadillac     381
Mazda        321
BMW          314
Infiniti     312
Suzuki       306
Audi         284
Mercedes-Benz 271
Hyundai      243
Kia          227
Subaru       210
Acura        210
Volvo        187
Lexus        186
Mitsubishi   175
Buick        166
Chrysler     159
Lincoln      152
Pontiac      141
Land Rover   123
Porsche      123
Aston Martin  91
Saab         78
Bentley      74
Ferrari      69
FIAT         62
Scion        60
Oldsmobile   57
Name: Make, dtype: int64
```

3.2. Define a `Tournament` class [30 marks](#)

You are to translate **these requirements** (with consideration of their usage, defined in **3.3** and **3.4**) into working code.

Overall: Coding efficiency and structure, including comments and docstrings, where appropriate, will contribute to the mark in the

- The class is initialised with:
 - a `DataFrame` of `cars`, as designed above
 - **No missing values** are permitted.
 - a tournament **name**
 - an **optional number of competing teams**, defaulting to 16.
 - If the input number of teams is not an integer, `raise` the appropriate kind of exception, and a message saying, "The number of teams must be an integer."
 - Also, `assert` that the value is positive and non-zero.
 - Ensure that this number is a **perfect square**.
 - Include sensible object representation dunder methods (i.e. `__repr__` and `__str__`).

- There is a method to `generate_sponsors`. - The method, by default, ****randomly selects a sponsor**** for each team from the available list of `Make`s in the `cars` `DataFrame`.
 - `generate_teams`. - This method should simply populate a list of `Team`'s in the `Tournament` class. - The teams are members of the `Team` class.

- There is a method to `buy_cars`.

- This method will allow the `Team`s to each purchase their initial inventory.

– There is a method to `_purchase_inventory`. – This internal method takes 1 argument : a single `Team` object. – With the information provided

- There is a method to `hold_event` (i.e., execute the tournament competition process).
- Cycle through the pairwise matches, keeping track of `Team` performance metrics.
 - The teams will compete in a head to head match and either continue on or be **eliminated**.
- After each match, allocate a **financial prize** to the winning `Team`. You can decide how to implement this; perhaps the prize increases in every
- After awarding the prize, allow the `Team` to `_purchase_inventory` again (increasing the number of cars in their inventory) before the next match.
 - Newly purchased cars can be duplicates of members of the `Team`'s existing inventory, but only one any kind of car is permitted to be purchased.
- At the end of the tournament event, record the `Tournament.champion` `Team`.

In [49]:

```
import pandas as pd
import numpy as np
from random import choice, randint, sample

class Tournament:
    """Tournement class"""
    class Team:
        def __init__(self, sponsor, budget):
            self.sponsor = sponsor
            self.budget = budget
            self.inventory = []
            self.active = True
            self.win_record = []

        def __str__(self):
            return f"Team sponsored by {self.sponsor} with budget £{self.budget} and {len(self.inventory)} cars."

    def __init__(self, cars, name, num_teams=16):
        if not isinstance(num_teams, int):
            raise ValueError("The number of teams must be an integer.")
        assert num_teams > 0, "Number of teams must be positive and non-zero."
        assert (num_teams & (num_teams - 1) == 0) and num_teams != 0, "Number of teams must be a perfect square."

        self.cars = cars
        self.name = name
        self.num_teams = num_teams
        self.sponsors = []
        self.budgets = []

    def __repr__(self):
        return f"Tournament(name={self.name}, num_teams={self.num_teams})"

    def __str__(self):
        return f"Tournament: {self.name} with {self.num_teams} teams competing."

    def generate_sponsors(self, specified_sponsors=None, low=100000, high=500000, incr=100):
        """
        Randomly selects sponsors for each team and random budgets between 100k and 500k.
        """
        makers = self.cars['Make'].unique()
        num_specified = len(specified_sponsors) if specified_sponsors else 0
        if num_specified < self.num_teams:
            random_sponsors = sample(list(makers), self.num_teams - num_specified)
            self.sponsors = specified_sponsors + random_sponsors if specified_sponsors else random_sponsors
        else:
            self.sponsors = specified_sponsors[:self.num_teams]

        self.budgets = [randint(low // incr, high // incr) * incr for _ in range(self.num_teams)]

    def generate_teams(self):
        """
        Generates the team objects within the tournament
        """
        self.teams = [self.Team(sponsor, budget) for sponsor, budget in zip(self.sponsors, self.budgets)]

    def buy_cars(self):
        """
        Method for teams to purchase their initial inventory based on their budget
        """
```

```

for team in self.teams:
    self._purchase_inventory(team)

def _purchase_inventory(self, team):
    """
    Internal method to buy cars from available cars
    """
    available_cars = self.cars[self.cars['Make'] == team.sponsor].copy()
    available_cars.sort_values(by='MPG-H', ascending=False, inplace=True)

    for _, car in available_cars.iterrows():
        if car['Price'] <= team.budget:
            team.inventory.append(car)
            team.budget -= car['Price']

def show_win_record(self):
    if self.teams is not None:
        for team in self.teams:
            print(f"{team.sponsor}: {team.win_record}")
    else:
        print("Teams not yet generated or competition not held.")

def show_teams(self):
    """
    Displays information about teams in the tournament including inventory size, budget and sponsor
    """
    if not self.teams:
        print("No teams have been generated yet.")
        return

    print(f"{'Team Sponsor':<20} {'Budget':<10} {'Inventory Size':<15}")
    for team in self.teams:
        sponsor = team.sponsor
        budget = f"${team.budget}"
        inventory_size = len(team.inventory)
        print(f"{'sponsor':<20} {'budget':<10} {'inventory_size':<15}")

def hold_event(self):
    """
    Executes the tournament and determines a champion
    """
    for team in self.teams:
        team.win_record = []

    round_teams = self.teams
    while len(round_teams) > 1:
        winners = []
        for i in range(0, len(round_teams), 2):
            team1, team2 = round_teams[i], round_teams[i+1]

            score1 = sum(car['MPG-H'] for car in team1.inventory)
            score2 = sum(car['MPG-H'] for car in team2.inventory)

            if score1 > score2:
                winner, loser = team1, team2
            else:
                winner, loser = team2, team1

            winner.win_record.append('W')
            loser.win_record.append('L')

            winners.append(winner)
        round_teams = winners

    self.champion = round_teams[0]

    print(f"The champion is {self.champion.sponsor} with a total MPG-H of {sum(car['MPG-H'] for car in self.champion.inventory)}")

```

3.3. Execute your `Tournament` class [5 marks](#)

- The process of building and executing the stages associated with the tournament will look like this:

```
t1 = Tournament(car, "The First Folks")
t1.generate_sponsors()
t1.generate_teams()
t1.buy_cars()
t1.hold_event()
print(f'The champion of {t1.name} Tournament is the {t1.champion}')
```

```
...
The champion of The First Folks Tournament is the Team sponsored by Nissan with $32000 available and 32 cars.
```

- Produce some visual (e.g., printed or plotted) record of the `Tournament` matches by invoking the `show_win_record` method:

```
t1.show_win_record()
```

```
...

Scion: ['W', 'W', 'L']
Suzuki: ['L']
Subaru: ['W', 'L']
Land Rover: ['L']
Nissan: ['W', 'W', 'W', 'W']
Dodge: ['L']
Lexus: ['L']
Saab: ['W', 'L']
Mercedes-Benz: ['L']
Buick: ['W', 'L']
Volkswagen: ['W', 'W', 'L']
Oldsmobile: ['L']
Hyundai: ['W', 'W', 'W', 'L']
Infiniti: ['L']
Cadillac: ['L']
BMW: ['W', 'L']
```

In [55]:

```
t1 = Tournament(cars, "SuperLeaded League")
t1.generate_sponsors()
t1.generate_teams()
t1.buy_cars()
t1.hold_event()
print(f'The champion of {t1.name} Tournament is the {t1.champion}')
t1.show_win_record()
```

```
The champion is Nissan with a total MPG-H of 1410
The champion of SuperLeaded League Tournament is the Team sponsored by Nissan with budget £1200 and 20 cars.
Lincoln: ['L']
Lexus: ['W', 'L']
Land Rover: ['L']
Infiniti: ['W', 'W', 'L']
GMC: ['L']
Kia: ['W', 'W', 'W', 'L']
BMW: ['L']
FIAT: ['W', 'L']
Volkswagen: ['L']
Mitsubishi: ['W', 'L']
Nissan: ['W', 'W', 'W', 'W']
Dodge: ['L']
Cadillac: ['L']
Chrysler: ['W', 'W', 'L']
Audi: ['W', 'L']
Saab: ['L']
```

In [56]:

```
t1.show_teams()
```

Team	Sponsor	Budget	Inventory Size
Lincoln		\$11800	3
Lexus		\$21900	5
Land Rover		\$2100	3
Infiniti		\$13600	9

GMC	\$1900	9
Kia	\$4700	10
BMW	\$3300	5
FIAT	\$800	11
Volkswagen	\$100	12
Mitsubishi	\$10100	25
Nissan	\$1200	20
Dodge	\$12000	25
Cadillac	\$24200	9
Chrysler	\$12200	18
Audi	\$7000	4
Saab	\$20100	10

2389

3.4. Compare the results of multiple `Tournament` executions [5 marks](#)

- Lastly, **execute 2** `Tournaments` in full.
- Compare the performance of their `.champion`s and produce a ranked representation of the different `Tournaments`. For example, when we rank tournaments by the score of their overall champion, we might produce:

Tournament ranking:

Position	Name	Sponsor	Score
1	The Other Group	Chevrolet	2044
2	The First Folks	Nissan	1737

In [57]:

```

tournament1 = Tournament(cars, name="The First Folks")
tournament1.generate_sponsors()
tournament1.generate_teams()
tournament1.buy_cars()
tournament1.hold_event()

tournament2 = Tournament(cars, name="SuperLeaded League")
tournament2.generate_sponsors()
tournament2.generate_teams()
tournament2.buy_cars()
tournament2.hold_event()

champion1_score = sum(car["MPG-H"] for car in tournament1.champion.inventory)
champion2_score = sum(car["MPG-H"] for car in tournament2.champion.inventory)

tournaments = [
    {"Position": 1, "Name": tournament1.name, "Sponsor": tournament1.champion.sponsor, "Score": champion1_score},
    {"Position": 2, "Name": tournament2.name, "Sponsor": tournament2.champion.sponsor, "Score": champion2_score}
]

tournaments_sorted = sorted(tournaments, key=lambda x: x["Score"], reverse=True)

for i, tournament in enumerate(tournaments_sorted, start=1):
    tournament["Position"] = i

print("Tournament ranking:")
print(f'{"Position":<12} {"Name":<25} {"Sponsor":<25} {"Score":<5}')
for tournament in tournaments_sorted:
    print(f'{"tournament["Position"]":<12} {"tournament["Name"]":<25} {"tournament["Sponsor"]":<25} {"tournament["Score"]":<5}')

```

The champion is Nissan with a total MPG-H of 1286

The champion is Toyota with a total MPG-H of 763

Tournament ranking:

Position	Name	Sponsor	Score
1	The First Folks	Nissan	1286
2	SuperLeaded League	Toyota	763

In []:

--

▼ cardata_modified.csv		 Download
1	Large file hidden. You can download it using the button above.	